

System Programming – Day 9

Rust

Nils Asmussen

10/16/2020

What is Rust?

*A language empowering everyone
to build reliable and efficient software.*

(rust-lang.org)

Why Another Language?

- We have plenty of languages to build reliably software:
 - Java, C#, Go, Python, Ruby, ...
 - All of these trade performance for safety
 - All of them have a runtime (garbage collector, ...)

Why Another Language?

- We have plenty of languages to build reliably software:
 - Java, C#, Go, Python, Ruby, ...
 - All of these trade performance for safety
 - All of them have a runtime (garbage collector, ...)
- We have plenty of languages to build efficient software:
 - C, C++, D, Assembly, ...
 - All of them trade safety for performance

Why Another Language?


- We have plenty of languages to build reliably software:
 - Java, C#, Go, Python, Ruby, ...
 - All of these trade performance for safety
 - All of them have a runtime (garbage collector, ...)
- We have plenty of languages to build efficient software:
 - C, C++, D, Assembly, ...
 - All of them trade safety for performance
- System programming requires efficiency/control *and* safety!

But Good Developers Don't Need Safety!


But Good Developers Don't Need Safety!

We closely study the root cause trends of vulnerabilities & search for patterns

% of memory safety vs. non-memory safety CVEs by patch year



But Good Developers Don't Need Safety!



General Idea of Rust

- C/C++ declare everything that is unsafe as “undefined behavior”
 - That pushes the problem to the developer
 - There is no way out: the developer has the control all the time

General Idea of Rust

- C/C++ declare everything that is unsafe as “undefined behavior”
 - That pushes the problem to the developer
 - There is no way out: the developer has the control all the time
- Rust provides safety without undefined behavior by default
 - The developer can opt out by marking code as “unsafe”
 - The developer *only* has the control if explicitly requested

General Idea of Rust

- C/C++ declare everything that is unsafe as “undefined behavior”
 - That pushes the problem to the developer
 - There is no way out: the developer has the control all the time
- Rust provides safety without undefined behavior by default
 - The developer can opt out by marking code as “unsafe”
 - The developer *only* has the control if explicitly requested
- Rust tracks ownership at compile time and thereby is
 - memory safe
 - data-race free

Agenda

Morning

- Getting started
- Ownership
- Basic Features + Exercise
- Structs and Enums + Exercise
- Generics, Traits, and Error Handling + Exercise

Agenda

Morning

- Getting started
- Ownership
- Basic Features + Exercise
- Structs and Enums + Exercise
- Generics, Traits, and Error Handling + Exercise

Afternoon

- Actual system programming with Rust
- You get a small kernel and implement semaphores

Repository

To get the slides and the exercises:

```
$ git clone https://tudos.org/~nils/sysprog-rust.git
```

Outline

- 1 Getting Started
- 2 Ownership
- 3 Basic Features
- 4 Structs, Enums, and Closures
- 5 Generics, Traits, and Error Handling
- 6 Advanced Features

Installation

- We need rustup (not rustc) to install the nightly version
- Some distributions (e.g., Arch) have a package for rustup
- Otherwise:

```
$ curl --proto '=https' --tlsv1.2 https://sh.rustup.rs -sSf > rup.sh  
# check if it's safe and use a fresh shell  
$ sh rup.sh # choose nightly, otherwise default values
```


Overview

- `rustc` is the Rust compiler; almost never invoked by the user
- `cargo` is Rust's build system and package manager
 - `Cargo.toml` describes what to build and its dependencies
 - `cargo` downloads dependencies and builds everything automatically
 - Every library/application is a *crate*
 - Crates can be found on <https://lib.rs> (or <https://crates.io>)

Let's Build Hello World!

```
$ cargo new hello
```

```
$ cd hello
```

```
$ cargo run
```

Outline

- 1 Getting Started
- 2 Ownership**
- 3 Basic Features
- 4 Structs, Enums, and Closures
- 5 Generics, Traits, and Error Handling
- 6 Advanced Features

Different Memory Management Approaches

- Many high-level languages use garbage collection to manage memory
 - Often not acceptable for OSes, bootloaders, VMMs, ...

Different Memory Management Approaches

- Many high-level languages use garbage collection to manage memory
 - Often not acceptable for OSes, bootloaders, VMMs, ...
- Many low-level languages let the developer manage memory explicitly
 - Error prone and the main cause for memory-safety issues

Different Memory Management Approaches

- Many high-level languages use garbage collection to manage memory
 - Often not acceptable for OSes, bootloaders, VMMs, ...
- Many low-level languages let the developer manage memory explicitly
 - Error prone and the main cause for memory-safety issues
- Rust uses Ownership
 - No garbage collection, no manual allocation
 - The compiler defines a set of rules and enforces them

Ownership Rules

- 1 Each value has a variable that's called its *owner*.
- 2 There can only be one owner at a time.
- 3 When the owner goes out of scope, the value will be *dropped*.

Ownership Rules – Examples

Valid example

```
{  
    let mut var = 4;    // mutable variable  
    var += 1;           // we are the owner  
} // var is dropped
```


Ownership Rules – Examples

Valid example

```
{  
    let mut var = 4;    // mutable variable  
    var += 1;          // we are the owner  
} // var is dropped
```

Invalid example

```
let mut var = 4;  
let var_ref = &mut var; // mutable reference to modify `var`  
drop(var);              // explicit drop  
*var_ref = 5;           // error (use after free)
```

Ownership Transfer and Borrowing

- 1 The owner of a value can *transfer* the ownership to someone else.

```
let var = String::from("hello"); // heap-allocated string
fn foo(name: String) { /* name is dropped */ }
foo(var);                  // transfer ownership to foo
```

Ownership Transfer and Borrowing

- 1 The owner of a value can *transfer* the ownership to someone else.

```
let var = String::from("hello"); // heap-allocated string
fn foo(name: String) { /* name is dropped */ }
foo(var);                  // transfer ownership to foo
```

- 2 Others can *borrow* a value from the owner.

```
let mut var = String::from("hello"); // mutable String
fn foo(name: &String) { /* use name */ }
foo(&var);                  // let foo borrow var
var.push(' ');              // we are the owner again
```

Outline

- 1 Getting Started
- 2 Ownership
- 3 Basic Features**
- 4 Structs, Enums, and Closures
- 5 Generics, Traits, and Error Handling
- 6 Advanced Features

Data Types (1)

- Scalars

- Integers: u32, i64, usize, . . .
- Floats: f32, f64
- Boolean: bool
- Character: char

- Structs

```
struct Foo {  
    field1: u32,  
    field2: String,  
}
```

Data Types (2)

- Tuples


```
let mut tuple = (1, "foo", 42);    // tuple length is fixed
tuple.0 += 1;                       // values are mutable
let (x, y, z) = tuple;              // destructuring
```

- Arrays

```
let mut array: [u32; 2] = [1, 2];  // arrays have a fixed size
array[3] += 1;                      // runtime error (bounds checked)
let foo = [0; 12];                  // array with 12 elements with value 0
```

Strings and Slices


```
let s = String::from("hello world");  
// String ~= Vec<char>  
let world = &s[6..11];  
// &str ~= &[char]
```



Strings and Slices

```
let s = String::from("hello world");  
// String ~= Vec<char>  
let world = &s[6..11];  
// &str ~= &[char]
```

```
&s[0..11]    // = "hello world"  
&s[6..]      // = "world"  
&s[..5]      // = "hello"  
&s[..]       // = "hello world"
```

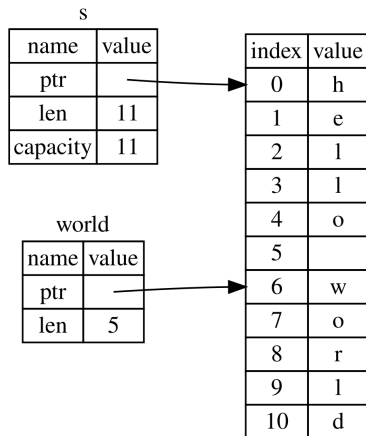


Strings and Slices

```
let s = String::from("hello world");  
// String ~= Vec<char>  
let world = &s[6..11];  
// &str ~= &[char]
```

```
&s[0..11]    // = "hello world"  
&s[6..]      // = "world"  
&s[..5]      // = "hello"  
&s[..]       // = "hello world"
```

```
let a = [1, 2, 3];  
&a[0..1]     // = [1]
```



Control Structures

- If expressions

```
if condition { println!("foo"); } else { println!("bar"); }  
let val = if condition { 4 } else { 5 };
```

- Loop

```
loop { }
```

- While

```
while condition { }
```

- For

```
for i in 0..10 { }
```

Exercise 1 – String Operations

- First exercise is in directory “words”
- Fill in the implementation of the functions
- Use the existing tests to verify your implementation:

```
$ cargo test
```

- Hint: use the standard library (<https://doc.rust-lang.org/stable>):
 - `str::chars`
 - `char::is_uppercase`
 - `str::split_whitespace`

Outline

- 1 Getting Started
- 2 Ownership
- 3 Basic Features
- 4 Structs, Enums, and Closures**
- 5 Generics, Traits, and Error Handling
- 6 Advanced Features

More on Structs

- Struct definition

```
struct Rectangle {  
    width: u32,  
    height: u32,  
}
```

More on Structs

- Struct definition

```
struct Rectangle {  
    width: u32,  
    height: u32,  
}
```

- Methods

```
impl Rectangle {  
    fn area(&self) -> u32 {  
        self.width * self.height  
    }  
}
```

More on Structs

- Struct definition

```
struct Rectangle {  
    width: u32,  
    height: u32,  
}
```

- Methods

```
impl Rectangle {  
    fn area(&self) -> u32 {  
        self.width * self.height  
    }  
}
```

- Method with mutable self

```
fn widen(&mut self, amount: u32) {  
    self.width += amount;  
}
```

More on Structs

- Struct definition

```
struct Rectangle {  
    width: u32,  
    height: u32,  
}
```

- Methods

```
impl Rectangle {  
    fn area(&self) -> u32 {  
        self.width * self.height  
    }  
}
```

- Method with mutable self

```
fn widen(&mut self, amount: u32) {  
    self.width += amount;  
}
```

- Method that take ownership

```
fn flip(self) -> Rectangle {  
    Rectangle {  
        width: self.height,  
        height: self.width,  
    }  
}
```


Enums

- Simple enumeration (like in C++)

```
enum Animal {  
    Sheep,  
    Cow,  
}
```

Enums

- Simple enumeration (like in C++)

```
enum Animal {  
    Sheep,  
    Cow,  
}
```

- Enums with data (tagged union)

```
enum Message {  
    Open(String),  
    Read(usize, usize),  
}
```

Enums

- Simple enumeration (like in C++)

```
enum Animal {  
    Sheep,  
    Cow,  
}
```

- Enums with data (tagged union)

```
enum Message {  
    Open(String),  
    Read(usize, usize),  
}
```

- Construction

```
Message::Open(String::from("Hello!"));  
Message::Read(0, 1024);
```

Enums

- Simple enumeration (like in C++)

```
enum Animal {  
    Sheep,  
    Cow,  
}
```

- Enums with data (tagged union)

```
enum Message {  
    Open(String),  
    Read(usize, usize),  
}
```

- Construction

```
Message::Open(String::from("Hello!"));  
Message::Read(0, 1024);
```

- Matching

```
match msg {  
    Message::Open(filename) => ...,  
    _ => println!("Unsupported"),  
}  
if let Message::Read(pos, num) = msg {  
}
```

Closure Basics

- Closures are anonymous functions that can be stored:

```
let adder = |x| { x += 1 };
```

Closure Basics

- Closures are anonymous functions that can be stored:

```
let adder = |x| { x += 1 };
```

- Closures can also capture their environment:

```
fn foo() {  
    let y = 42;  
    let adder = |x| { x += y };  
}
```

Closure Representations

- 1 Fn: capture environment by immutable references
- 2 FnMut: capture environment by mutable references
- 3 FnOnce: capture environment by ownership transfer

Closure Representations

- 1 Fn: capture environment by immutable references
- 2 FnMut: capture environment by mutable references
- 3 FnOnce: capture environment by ownership transfer

Example

```
fn count<F: Fn(&u32) -> bool>(vec: &[u32], func: F) -> usize {  
    let mut count = 0;  
    for e in vec {  
        if func(e) { count += 1; }  
    }  
    count  
}
```


Exercise 2 – Command Line People Database

- Second exercise is in directory “people”
- Simple command line program that lets the user manage a “database” of people
- Fill in the missing parts (parsing, command execution)
- It’s okay to use `unwrap/panic` (we’ll add proper error handling later)
- The following building blocks might be helpful:
 - `Iterator::collect`
 - `Iterator::find`
 - `Vec::push`
 - `Vec::retain`

Outline

- 1 Getting Started
- 2 Ownership
- 3 Basic Features
- 4 Structs, Enums, and Closures
- 5 Generics, Traits, and Error Handling**
- 6 Advanced Features

Basics of Generics

- Generics allow to define functions/structs/enums for a variety of concrete types:

```
fn foo<T>(arg: T) { /* ... */ }
```

- Generics have no runtime overhead due to monomorphization:

```
fn foo<T>(arg: T) { /* ... */ }
```

// is compiled to something like:

```
fn foo_u32(arg: u32) { /* ... */ }
```

```
fn foo_u64(arg: u64) { /* ... */ }
```

- Unlike C++, Rust is strict about the requirements for type parameters (based on traits, as we will see shortly)

Generic Types

- Generic function

```
fn head<T>(elems: &Vec<T>) -> &T {  
    &elems[0]  
}  
assert_eq!(*head(&vec![1, 2]), 1);
```

Generic Types

- Generic function

```
fn head<T>(elems: &Vec<T>) -> &T {  
    &elems[0]  
}  
assert_eq!(*head(&vec![1, 2]), 1);
```

- Generic struct

```
struct Rectangle<T> {  
    width: T,  
    height: T,  
}  
Rectangle { width: 1.2, height: 4.5 }
```

Generic Types

- Generic function

```
fn head<T>(elems: &Vec<T>) -> &T {  
    &elems[0]  
}  
assert_eq!(*head(&vec![1, 2]), 1);
```

- Generic struct

```
struct Rectangle<T> {  
    width: T,  
    height: T,  
}  
Rectangle { width: 1.2, height: 4.5 }
```

- Generic enum

```
enum Option<T> {  
    Some(T),  
    None,  
}
```

Generic Types

- Generic function

```
fn head<T>(elems: &Vec<T>) -> &T {  
    &elems[0]  
}  
assert_eq!(*head(&vec![1, 2]), 1);
```

- Generic struct

```
struct Rectangle<T> {  
    width: T,  
    height: T,  
}  
Rectangle { width: 1.2, height: 4.5 }
```

- Generic enum

```
enum Option<T> {  
    Some(T),  
    None,  
}
```

- Generic method

```
impl<T: AddAssign> Rectangle<T> {  
    fn widen(&mut self, amount: T) {  
        self.width += amount;  
    }  
}
```

Trait Basics

- A *trait* defines a behavior that can be implemented by multiple types:

```
trait Shape {  
    fn area(&self) -> u32;  
}
```


Trait Basics

- A *trait* defines a behavior that can be implemented by multiple types:

```
trait Shape {  
    fn area(&self) -> u32;  
}
```

- Implementing a trait for a type:

```
impl Shape for Rectangle {  
    fn area(&self) -> u32 {  
        self.width * self.height  
    }  
}
```

More on Traits (1)

- Using trait bounds:

```
fn sum<T: AddAssign + Copy + Default>(nums: &Vec<T>) -> T {  
    let mut sum = T::default();  
    for n in nums { sum += *n; }  
    sum  
}
```

More on Traits (1)

- Using trait bounds:

```
fn sum<T: AddAssign + Copy + Default>(nums: &Vec<T>) -> T {  
    let mut sum = T::default();  
    for n in nums { sum += *n; }  
    sum  
}
```

- Static vs. dynamic dispatch:

```
// one function for each type  
fn static_dispatch<T: Shape>(sh: &T) { }  
fn static_dispatch(sh: &impl Shape) { } // syntactic sugar  
// one function for all types, dispatched at runtime  
fn dynamic_dispatch(sh: &dyn Shape) { }
```

More on Traits (2)

- Derive attribute:

```
#[derive(Debug, Copy, Clone)]
```

```
struct Point {
```

```
    x: u32,
```

```
    y: u32,
```

```
}
```

More on Traits (2)

- Derive attribute:

```
#[derive(Debug, Copy, Clone)]
```

```
struct Point {
```

```
    x: u32,
```

```
    y: u32,
```

```
}
```

```
let p1 = Point { x: 0, y: 16 };
```

```
let p2 = p1;           // copied
```

```
println!("p = {:?}", p1); // prints "p = Point { x: 0, y: 16 }"
```

Error Handling

- Unrecoverable errors with panic!:
 - Sometimes the best you can do
 - Can perform stack unwinding or not (set panic=abort)
 - Provides a backtrace to the user

Error Handling

- Unrecoverable errors with `panic!`:
 - Sometimes the best you can do
 - Can perform stack unwinding or not (set `panic=abort`)
 - Provides a backtrace to the user
- Recoverable errors with `Result`:

```
enum Result<T, E> {  
    Ok(T),  
    Err(E),  
}
```

How to Handle Errors?

```
fn read_file(path: &str) -> Result<String, io::Error> {  
    let mut file = std::fs::File::open(path)?;  
    let mut s = String::new();  
    file.read_to_string(&mut s)?;  
    Ok(s)  
}  
let content = read_file("my_file.txt").expect("Unable to read file");
```


How to Handle Errors?

```
fn read_file(path: &str) -> Result<String, io::Error> {  
    let mut file = std::fs::File::open(path)?;  
    let mut s = String::new();  
    file.read_to_string(&mut s)?;  
    Ok(s)  
}  
  
let content = read_file("my_file.txt").expect("Unable to read file");  
  
let mut file = std::fs::File::open(path)?;  
// is equivalent to:  
let mut file = match std::fs::File::open(path) {  
    Ok(file) => file,  
    Err(e) => return Err(e),  
};
```

Option Instead of Nullpointers

- Similar to Result for errors, Rust uses Option for optional values:

```
enum Option<T> {  
    Some(T),  
    None,  
}
```

Option Instead of Nullpointers

- Similar to Result for errors, Rust uses Option for optional values:

```
enum Option<T> {  
    Some(T),  
    None,  
}
```

- Important methods on Result and Option
 - unwrap: panic if None/Err
 - expect: panic with message if None/Err
 - *_or_else: transformation

Option Instead of Nullpointers

- Similar to Result for errors, Rust uses Option for optional values:

```
enum Option<T> {  
    Some(T),  
    None,  
}
```

- Important methods on Result and Option
 - unwrap: panic if None/Err
 - expect: panic with message if None/Err
 - *_or_else: transformation
- More at <https://doc.rust-lang.org/stable>

Exercise 3 – Proper Error Handling

- Let's add proper error handling to our people “database”
- Use `Result` and `Option` where appropriate
- Get rid of all panics/unwraps
- Hints:
 - Introduce your own error enum
 - Attach `#[derive(Debug)]` to your error enum
 - Implement `From<std::num::ParseIntError>` for your enum
 - Implement `Display` for `Person`

Outline

- 1 Getting Started
- 2 Ownership
- 3 Basic Features
- 4 Structs, Enums, and Closures
- 5 Generics, Traits, and Error Handling
- 6 Advanced Features**

Unsafe

- Rust allows you to enable additional features via `unsafe`
- Tells the compiler that you know what you're doing
- Does *not* turn off safety checks, but allows you additionally to:
 - Dereference raw pointers
 - Call unsafe functions
 - (Some more that are less important)
- Unsafe code is typically used to build safe abstractions (`Vec`, `String`, ...)
- Example:

```
let mut_ptr = 0xB8000 as *mut u32;           // VGA frame buffer
let const_ptr = 0xDEAD_BEEF as *const u32;
unsafe { *mut_ptr = *const_ptr; }
```

Interfacing with Other Languages

- Rust can interface with other languages through the foreign function interface (FFI)
- Allows to call C functions from Rust:

```
extern "C" {  
    fn abs(input: i32) -> i32;  
}  
unsafe { abs(-2) };
```


Interfacing with Other Languages

- Rust can interface with other languages through the foreign function interface (FFI)
- Allows to call C functions from Rust:

```
extern "C" {  
    fn abs(input: i32) -> i32;  
}  
unsafe { abs(-2) };
```

- And to export Rust functions to C:

```
#[no_mangle]  
extern "C" fn rust_double(arg: u64) -> u64 {  
    arg * 2  
}
```

Interior Mutability

- Design pattern that allows to bypass the ownership model
- Used by Cell, RefCell, Mutex and others

Interior Mutability

- Design pattern that allows to bypass the ownership model
- Used by Cell, RefCell, Mutex and others
- For example (from the kernel you'll work on):

```
pub struct StaticCell<T: Sized> { inner: UnsafeCell<T> }
```

```
impl<T: Sized> StaticCell<T> {  
    pub const fn new(val: T) -> Self {  
        StaticCell { inner: UnsafeCell::new(val) }  
    }  
    pub fn get_mut(&self) -> &mut T {  
        unsafe { &mut *self.inner.get() }  
    }  
}
```

Exercise 4 – Semaphores

- Last exercise is in directory “kernel”
- Simple kernel that supports exactly two programs and runs in physical memory
- The program is instantiated two times and performs prints in a loop
- The prints currently mix occasionally; use semaphores to prevent it
- You need:
 - Add block support for tasks (preemptive scheduling already implemented)
 - Add the Semaphore itself (with up and down based on blockable tasks)
 - Add up/down syscalls to the kernel
 - Use the syscalls in the user program